



# Overview of ROCm and Compilers

**Samuel Antao**  
**STMS, LUMI CoE Lead**  
**Data Center GPU, AMD**

AMD @HLRS

**AMD**   
together we advance\_

# Thanks to all the AMD staff for their contributions

Suyash Tandon

Justin Chang

Julio Maia

Noel Chalmers

Paul T. Bauman

Nicholas Curtis

Nicholas Malaya

Alessandro Fanfarillo

Jose Noudohouenou

Chip Freitag

Damon McDougall

Noah Wolfe

Jakub Kurzak

Samuel Antao

George Markomanolis

Bob Robey

Gina Sitaraman

---

# Agenda

- 
1. ROCm Software Ecosystem
  2. Compilers for AMD GPUs

---

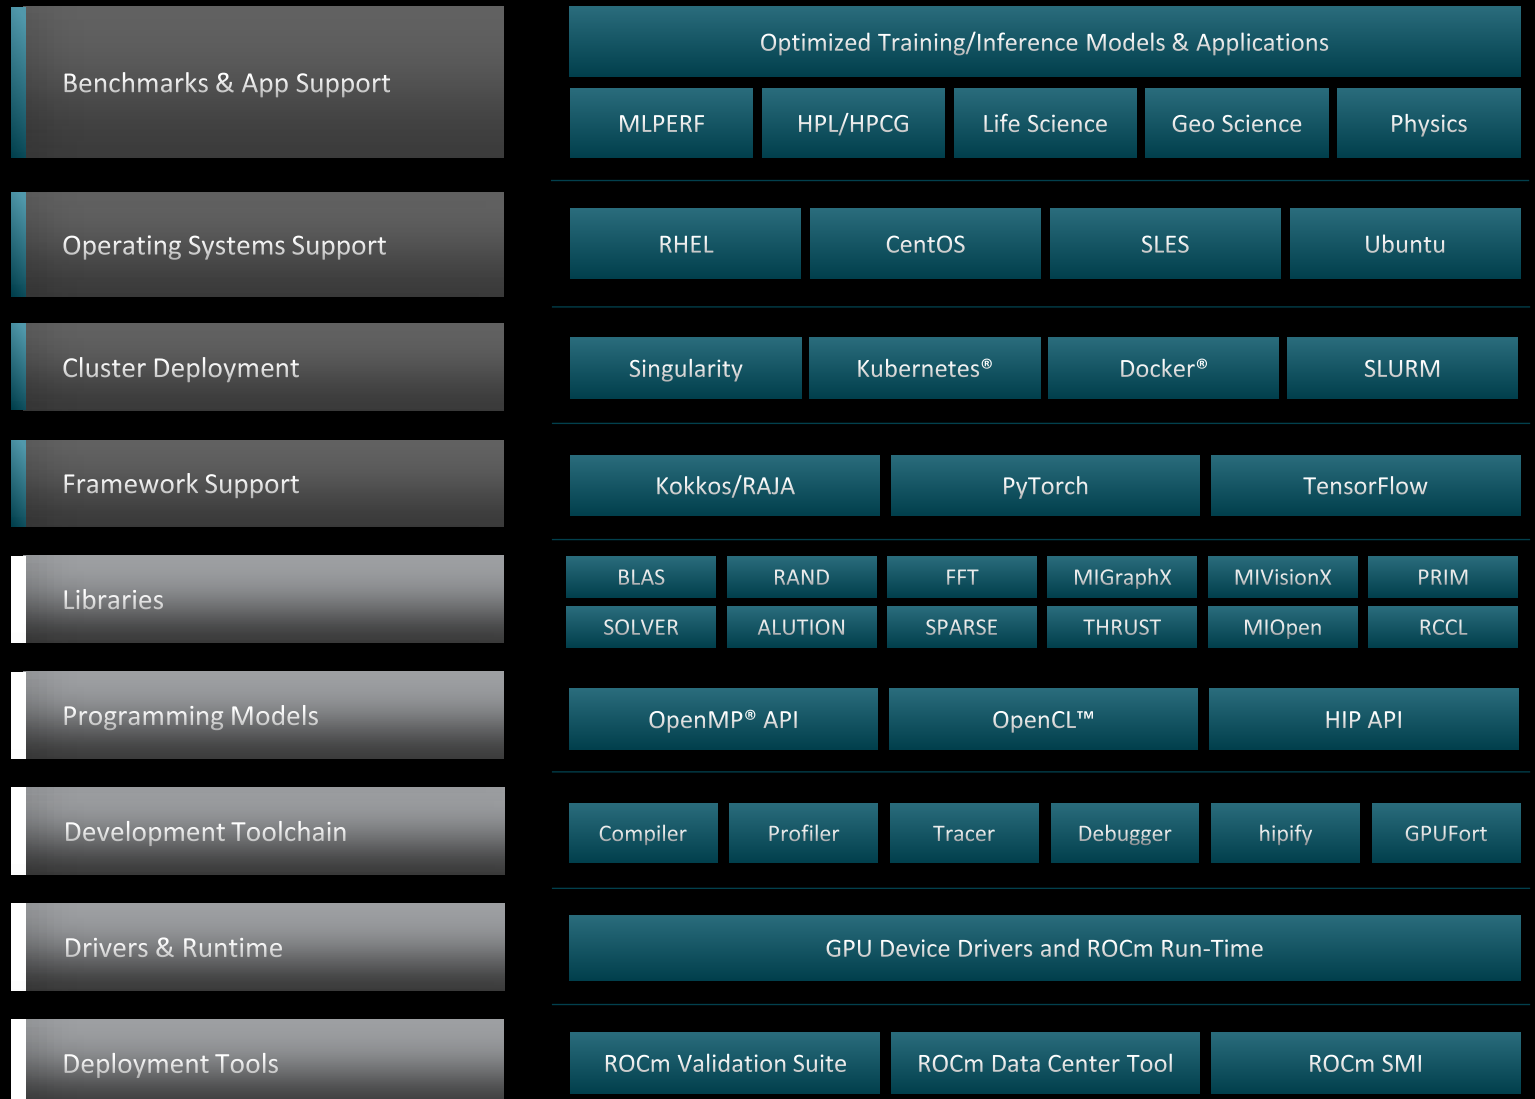
## 2. ROCm Software Ecosystem

---

# Open Software Platform For GPU Compute



- Unlocked GPU Power To Accelerate Computational Tasks
- Optimized for HPC and Deep Learning Workloads at Scale
- Open Source Enabling Innovation, Differentiation, and Collaboration



# AMD ROCm 5.0

DEMOCRATIZING EXASCALE FOR ALL

## EXPANDING SUPPORT & ACCESS

- Support for Radeon Pro W6800 Workstation GPUs
- Remote access through the AMD Accelerator Cloud

## OPTIMIZING PERFORMANCE

- MI200 Optimizations: FP64 Matrix ops, Improved Cache
- Improved launch latency and kernel performance

## ENABLING DEVELOPER SUCCESS

- HPC Apps & ML Frameworks on AMD InfinityHub
- Streamlined and improved tools increasing productivity

# ROCm Software Ecosystem

- **H**eterogeneous-compute **I**nterface for **P**ortability (HIP) is part of a larger software distribution called ROCm
- Install instructions and documentation:
  - [https://rocm.docs.amd.com/en/latest/deploy/linux/quick\\_start.html](https://rocm.docs.amd.com/en/latest/deploy/linux/quick_start.html)
  - <https://gpuopen.com/learn/amd-lab-notes/amd-lab-notes-rocm-installation-readme/>
- The ROCm package provides libraries and programming tools for developing HPC and ML applications on AMD GPUs
- All the ROCm environment and the libraries are provided from the supercomputer, usually, there is no need to install something yourselves
- Heterogeneous System Architecture (HSA) runtime is an API that exposes the necessary interfaces to access and interact with the hardware driven by AMDGPU driver



# ROCm GPU Libraries

ROCm provides several GPU math libraries

- Typically, two versions:
  - roc\* -> AMD GPU library, usually written in HIP
  - hip\* -> Thin interface between roc\* and Nvidia cu\* library

When developing an application meant to target both CUDA and AMD devices, use the hip\* libraries (portability)

When developing an application meant to target only AMD devices, may prefer the roc\* library API (performance).

- Some roc\* libraries perform **better** by using addition APIs not available in the cu\* equivalents

hipBLAS

rocBLAS

cuBLAS



# AMD Math Library Equivalents: “Decoder Ring”

<b>CUBLAS</b>	<b>ROCBLAS</b>	Basic Linear Algebra Subroutines
<b>CUFFT</b>	<b>ROCFFT</b>	Fast Fourier Transforms
<b>CURAND</b>	<b>ROCRAND</b>	Random Number Generation
<b>THRUST</b>	<b>ROCTHRUST</b>	C++ Parallel Algorithms
<b>CUB</b>	<b>ROCPRIM</b>	Optimized Parallel Primitives

# AMD Math Library Equivalents: “Decoder Ring”

**CUSPARSE**

**ROCSPARSE**

Sparse BLAS, SpMV, etc.

**CUSOLVER**

**ROCSOLVER**

Linear Solvers

**AMGX**


**ROCALUTION**

Solvers and preconditioners  
for sparse linear systems

[GITHUB.COM/ROCM-DEVELOPER-TOOLS/HIP](https://github.com/ROCm-developer-tools/HIP) → [HIP\\_PORTING\\_GUIDE.MD](#) FOR A COMPLETE LIST

# AMD GPU Libraries: BLAS

- rocBLAS – `sudo apt install rocblas`
  - Source code: <https://github.com/ROCmSoftwarePlatform/rocBLAS>
  - Documentation: <https://rocblas.readthedocs.io/en/latest/>
  - Basic linear algebra functionality
    - axpy, gemv, trsm, etc
  - Use hipBLAS if you need portability between AMD and NVIDIA devices
- hipBLAS - `sudo apt install hipblas`
  - Documentation: <https://github.com/ROCmSoftwarePlatform/hipBLAS/wiki/Exported-functions>
  - Use this if you need portability between AMD and NVIDIA
  - It is just a thin wrapper:
    - It can dispatch calls to rocBLAS for AMD devices
    - It can dispatch calls to cuBLAS for NVIDIA devices



# AMD GPU Libraries: rocBLAS example

- rocBLAS
  - Documentation: <https://rocblas.readthedocs.io/en/latest/>
  - Level 1, 2, and 3 functionality
    - axpy, gemv, trsm, etc
  - Note: rocBLAS syntax matches BLAS closer than hipBLAS or cuBLAS
    - Use hipBLAS only if you need portability between AMD and NVIDIA devices
  - Link with: `-lrocblas`

```
#include <rocblas.h>

int main(int argc, char ** argv) {
    rocblas_int N = 500000;

    // Allocate device memory
    double * dx, * dy;
    hipMalloc(&dx, sizeof(double) * N);
    hipMalloc(&dy, sizeof(double) * N);

    // Allocate host memory (and fill up the arrays) here
    std::vector<double> hx(N), hy(N);

    // Copy host arrays to device
    hipMemcpy(dx, hx.data(), sizeof(double) * N, hipMemcpyHostToDevice);
    hipMemcpy(dy, hy.data(), sizeof(double) * N, hipMemcpyHostToDevice);

    const double alpha = 1.0;
    rocblas_handle handle;
    rocblas_create_handle(&handle);
    rocblas_status status;
    status = rocblas_daxpy(handle, N, &alpha, dx, 1, dy, 1);
    rocblas_destroy_handle(handle);

    // Copy result back to host
    hipMemcpy(hy.data(), dy, sizeof(double) * N, hipMemcpyDeviceToHost);
    hipFree(dx);
    hipFree(dy);
    return 0;
}
```

AMD @HLRS

# Some Links to Key Libraries

- BLAS
  - rocBLAS (<https://github.com/ROCmSoftwarePlatform/rocBLAS>)
  - hipBLAS (<https://github.com/ROCmSoftwarePlatform/hipBLAS>)
- FFTs
  - rocFFT (<https://github.com/ROCmSoftwarePlatform/rocFFT>)
  - hipFFT (<https://github.com/ROCmSoftwarePlatform/hipFFT>)
- Random number generation
  - rocRAND (<https://github.com/ROCmSoftwarePlatform/rocRAND>)
- Sparse linear algebra
  - rocSPARSE (<https://github.com/ROCmSoftwarePlatform/rocSPARSE>)
  - hipSPARSE (<https://github.com/ROCmSoftwarePlatform/hipSPARSE>)
- Iterative solvers
  - rocALUTION (<https://github.com/ROCmSoftwarePlatform/rocALUTION>)
- Parallel primitives
  - rocPRIM (<https://github.com/ROCmSoftwarePlatform/rocPRIM>)
  - hipCUB (<https://github.com/ROCmSoftwarePlatform/hipCUB>)

# AMD Machine Learning Library Support

## Machine Learning Frameworks:

- Tensorflow: <https://github.com/ROCmSoftwarePlatform/tensorflow-upstream>
- Pytorch: <https://github.com/ROCmSoftwarePlatform/pytorch>
- Caffe: <https://github.com/ROCmSoftwarePlatform/hipCaffe>

## Machine Learning Libraries:

- MIOpen (similar to cuDNN): <https://github.com/ROCmSoftwarePlatform/MIOpen>
- Tensile (GEMM Autotuner): <https://github.com/ROCmSoftwarePlatform/Tensile>
- RCCL (ROCm analogue of NCCL): <https://github.com/ROCmSoftwarePlatform/rccl>
- Horovod (Distributed ML): <https://github.com/ROCmSoftwarePlatform/horovod>

## Benchmarks:

- DeepBench: <https://github.com/ROCmSoftwarePlatform/DeepBench>
- MLPerf: <https://mlperf.org>

# Usage of hipcc

Usage is straightforward. Accepts all/any flags that clang accepts, e.g.,

```
hipcc --offload-arch=gfx90a dotprod.cpp -o dotprod
```

Set HIPCC\_VERBOSE=7 to see a bunch of useful information

- Compile and link lines
- Various paths

```
$ HIPCC_VERBOSE=7 hipcc --offload-arch=gfx90a dotprod.cpp -o dotprod
```

```
HIP_PATH=/opt/rocm-5.2.0
```

```
HIP_PLATFORM=amd
```

```
HIP_COMPILER=clang
```

```
HIP_RUNTIME=rocc1r
```

```
ROCM_PATH=/opt/rocm-5.2.0
```

```
...
```

```
hipcc-args: --offload-arch=gfx90a dotprod.cpp -o dotprod
```

```
hipcc-cmd: /opt/rocm-5.2.0/llvm/bin/clang++ -stdc=c++11 -hc -D__HIPCC__ -isystem /opt/rocm-
```

```
5.2.0/llvm/lib/clang/14.0.0/include
```

```
-isystem /opt/rocm-5.2.0/has/include -isystem /opt/rocm-5.2.0/include -offload-arch=gfx90a -O3 ...
```

- You can use also *hipcc -v ...* to print some information
- With the command *hipconfig* you can see many information about environment variables declaration

# OpenMP Offload GPU Support

- ROCm and AOMP
  - ROCm supports both HIP and OpenMP
  - AOMP: the AMD OpenMP research compiler, it is used to prototype the new OpenMP features for ROCm
- HPE Compilers
  - Provides offloading support to AMD GPUs, through OpenMP, HIP, and OpenACC (only for Fortran)
- GNU compilers:
  - Provide OpenMP and OpenACC offloading support for AMD GPUs
  - GCC 11: Supports AMD GCN gfx908
  - GCC 13: Supports AMD GCN gfx90a



# Disclaimer

The information presented in this document is for informational purposes only and may contain technical inaccuracies, omissions, and typographical errors. The information contained herein is subject to change and may be rendered inaccurate for many reasons, including but not limited to product and roadmap changes, component and motherboard version changes, new model and/or product releases, product differences between differing manufacturers, software changes, BIOS flashes, firmware upgrades, or the like. Any computer system has risks of security vulnerabilities that cannot be completely prevented or mitigated. AMD assumes no obligation to update or otherwise correct or revise this information. However, AMD reserves the right to revise this information and to make changes from time to time to the content hereof without obligation of AMD to notify any person of such revisions or changes.

THIS INFORMATION IS PROVIDED ‘AS IS.’ AMD MAKES NO REPRESENTATIONS OR WARRANTIES WITH RESPECT TO THE CONTENTS HEREOF AND ASSUMES NO RESPONSIBILITY FOR ANY INACCURACIES, ERRORS, OR OMISSIONS THAT MAY APPEAR IN THIS INFORMATION. AMD SPECIFICALLY DISCLAIMS ANY IMPLIED WARRANTIES OF NON-INFRINGEMENT, MERCHANTABILITY, OR FITNESS FOR ANY PARTICULAR PURPOSE. IN NO EVENT WILL AMD BE LIABLE TO ANY PERSON FOR ANY RELIANCE, DIRECT, INDIRECT, SPECIAL, OR OTHER CONSEQUENTIAL DAMAGES ARISING FROM THE USE OF ANY INFORMATION CONTAINED HEREIN, EVEN IF AMD IS EXPRESSLY ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

Third-party content is licensed to you directly by the third party that owns the content and is not licensed to you by AMD. ALL LINKED THIRD-PARTY CONTENT IS PROVIDED “AS IS” WITHOUT A WARRANTY OF ANY KIND. USE OF SUCH THIRD-PARTY CONTENT IS DONE AT YOUR SOLE DISCRETION AND UNDER NO CIRCUMSTANCES WILL AMD BE LIABLE TO YOU FOR ANY THIRD-PARTY CONTENT. YOU ASSUME ALL RISK AND ARE SOLELY RESPONSIBLE FOR ANY DAMAGES THAT MAY ARISE FROM YOUR USE OF THIRD-PARTY CONTENT.

© 2023 Advanced Micro Devices, Inc. All rights reserved. AMD, the AMD Arrow logo, AMD CDNA, AMD ROCm, AMD Instinct, and combinations thereof are trademarks of Advanced Micro Devices, Inc. in the United States and/or other jurisdictions. Other names are for informational purposes only and may be trademarks of their respective owners.

Docker and the Docker logo are trademarks or registered trademarks of Docker, Inc.

Kubernetes is a registered trademark of The Linux® Foundation.

Linux is the registered trademark of Linus Torvalds in the U.S. and other countries.

OpenCL is a trademark of Apple Inc. used by permission by Khronos Group, Inc.

The OpenMP name and the OpenMP logo are registered trademarks of the OpenMP Architecture Review Board

# Questions?

